



# Recommender Systems


Shantanu Jain



# Recommender Systems

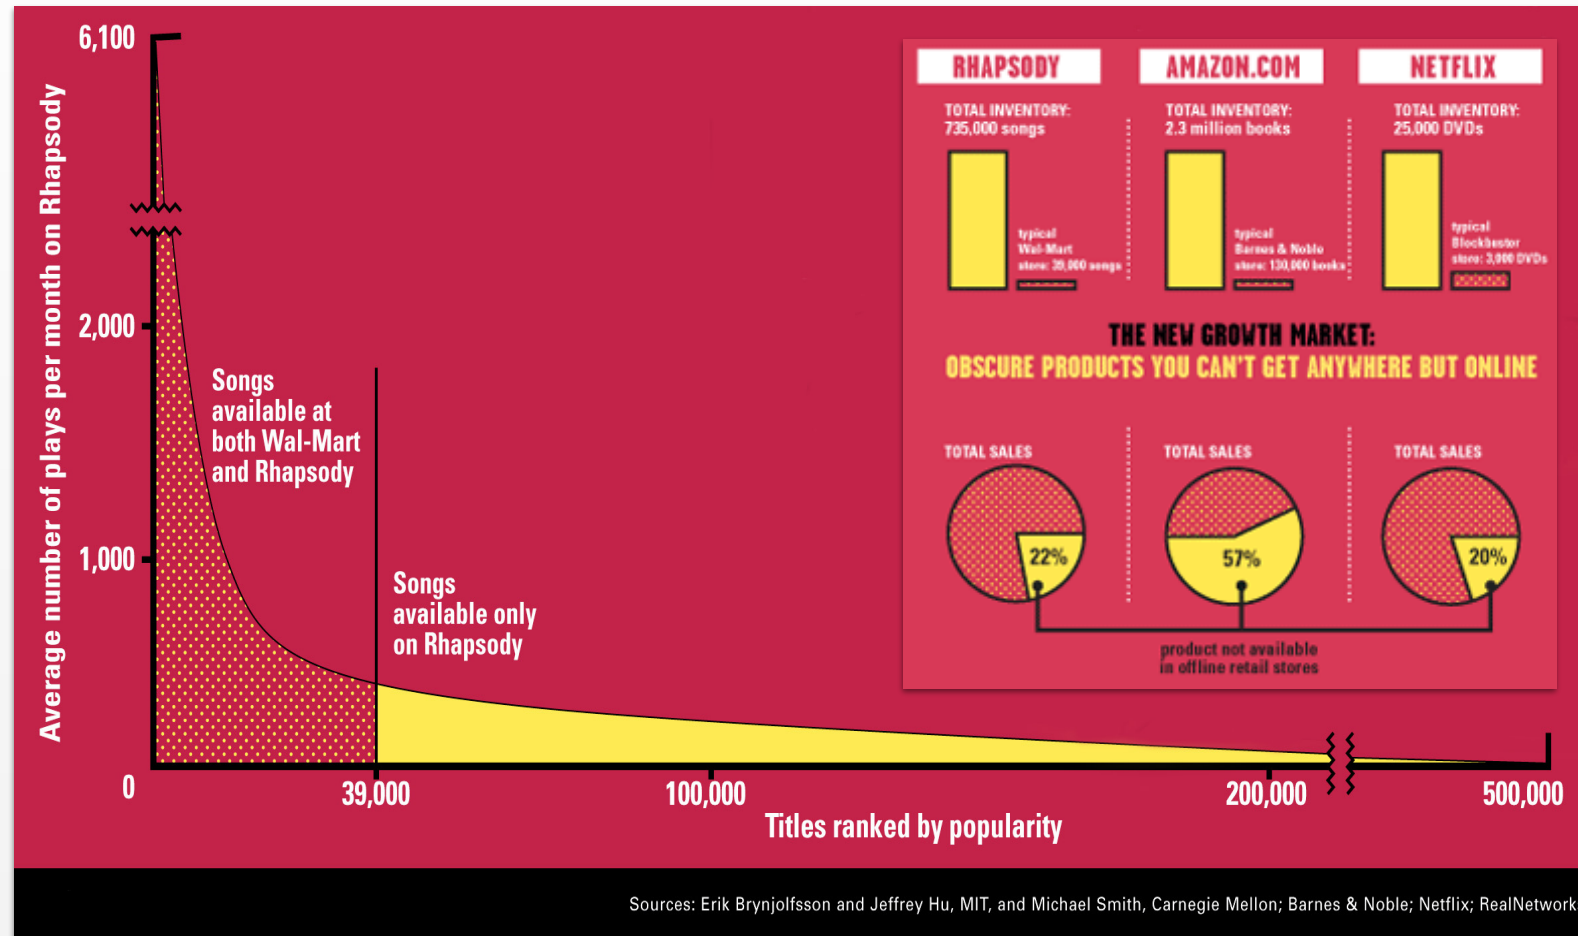
Reasoning about the Long Tail

# The Long Tail



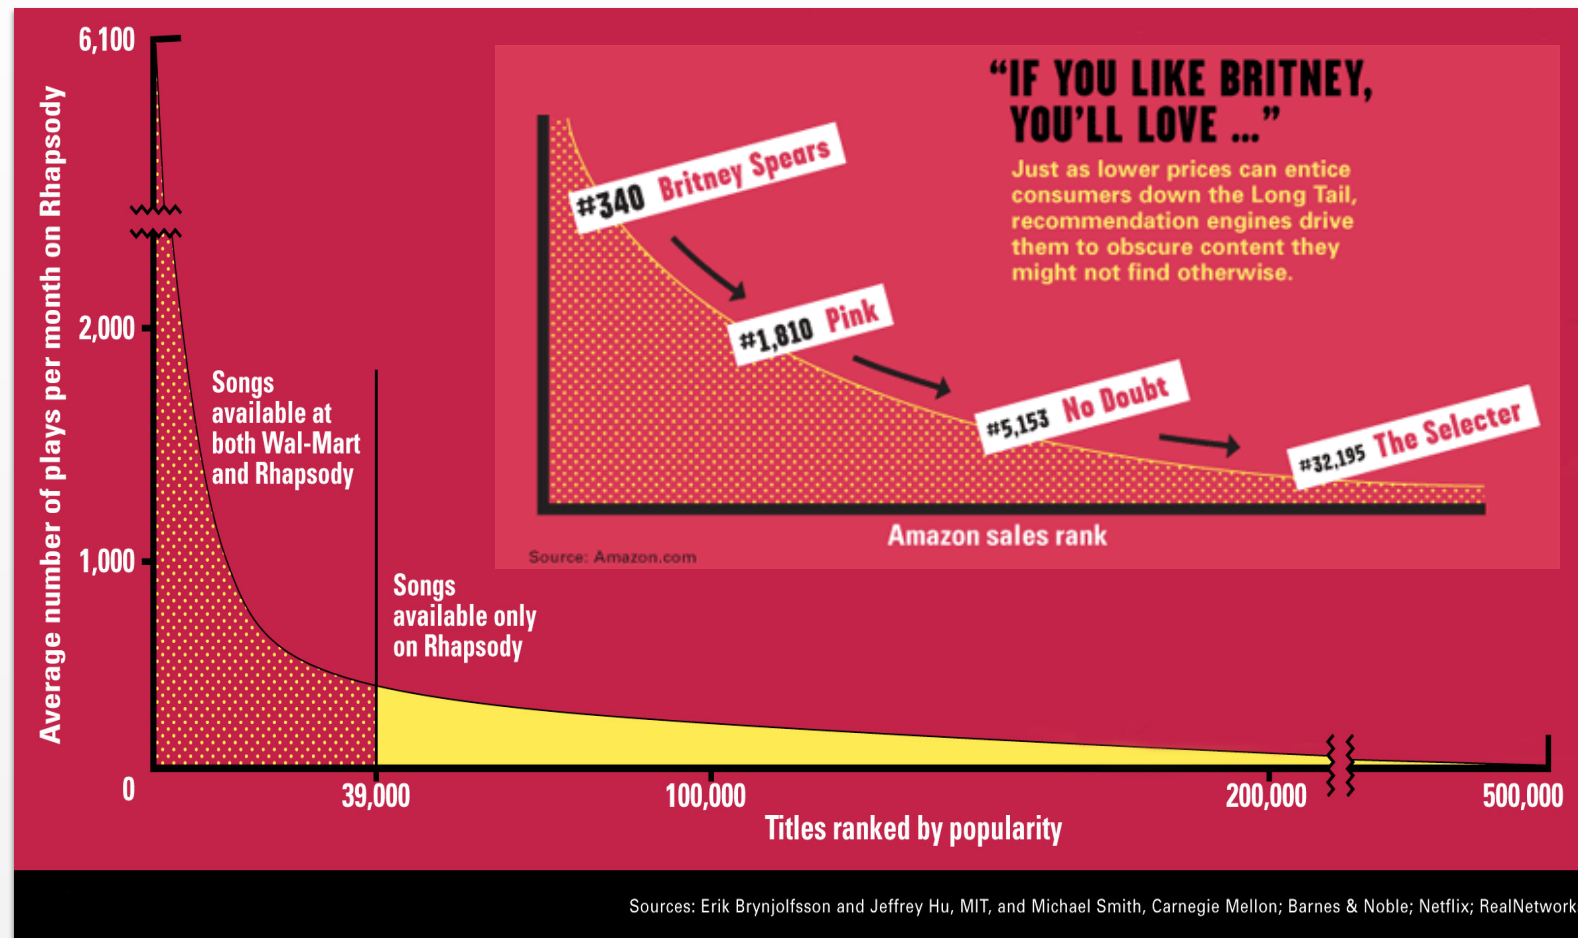
(from: <https://www.wired.com/2004/10/tail/>)

# The Long Tail



(from: <https://www.wired.com/2004/10/tail/>)

# The Long Tail



(from: <https://www.wired.com/2004/10/tail/>)

# Applications of Recommender Systems

- Movie recommendation (Netflix)
- Related product recommendation (Amazon)
- Web page ranking (Google)
- Social recommendation (Facebook)
- Priority inbox & spam filtering (Google)
- Online dating (OK Cupid)
- Computational Advertising (Everyone)

# Problem Setting

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

# Problem Setting

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?



# Problem Setting

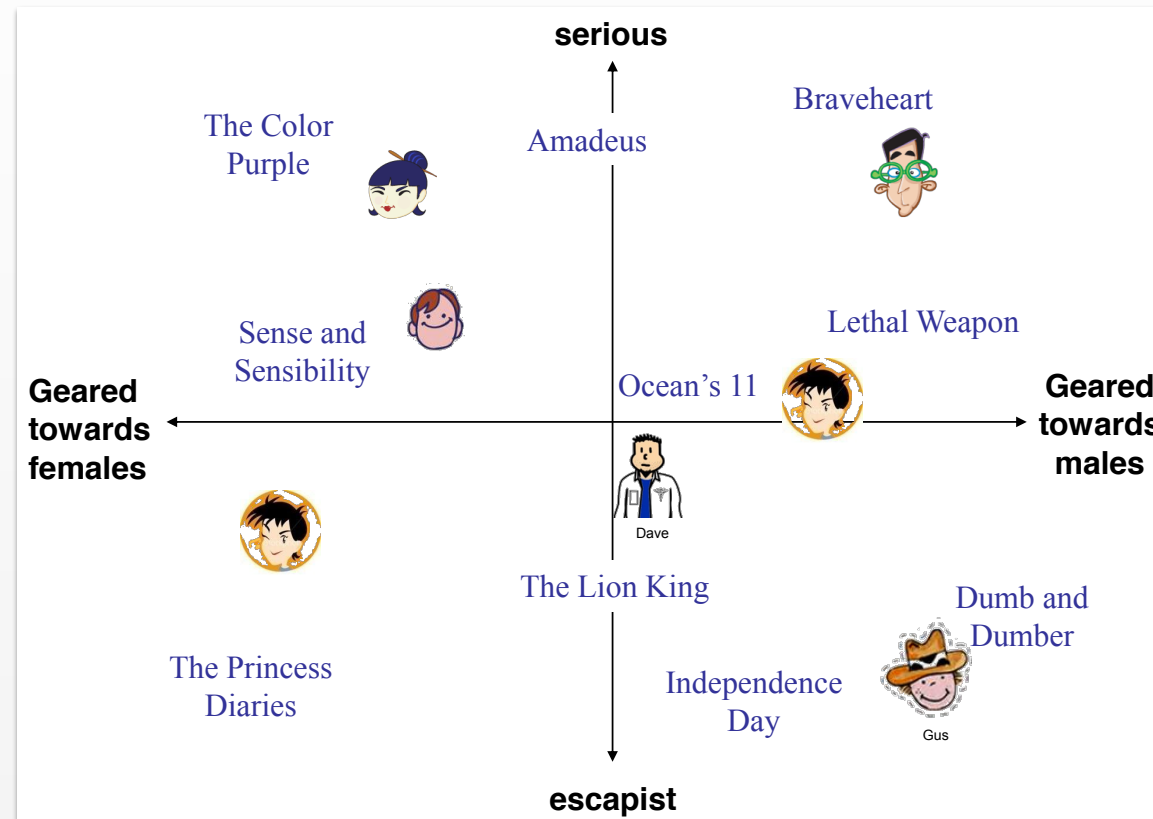
Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

# Problem Setting

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

- *Task*: Predict user preferences for unseen items

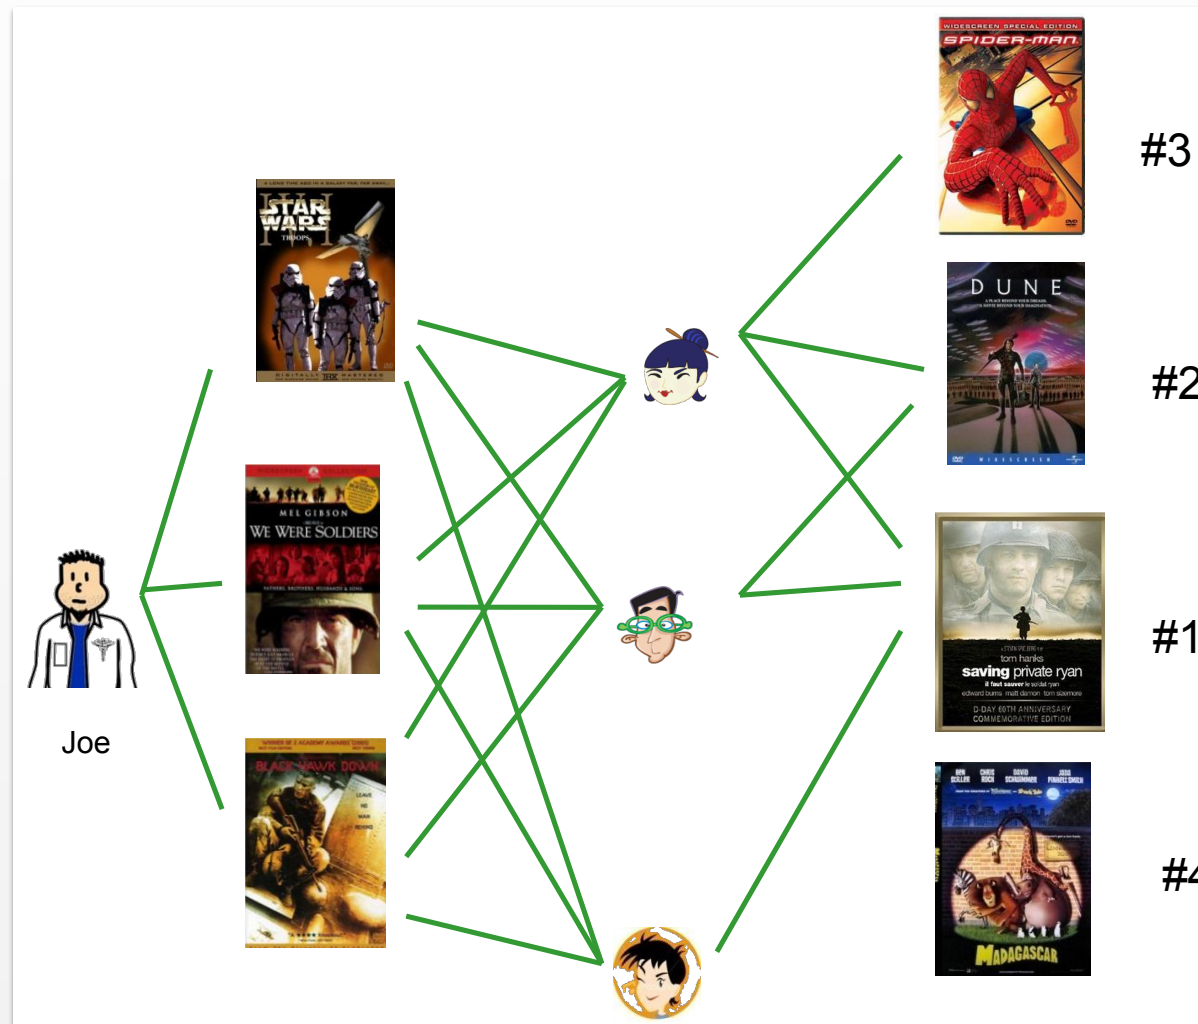
# Content-based Filtering



*Two Approaches:*

1. Predict rating using **item** features on a **per-user** basis
2. Predict rating using **user** features on a **per-item** basis

# Collaborative Filtering



*Idea:* Predict rating based on similarity to other users

# Problem Setting

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

- *Task*: Predict user preferences for unseen items
- *Content-based filtering*: Model user/item features
- *Collaborative filtering*: Implicit similarity of users or items

# Running Yardstick: RMSE

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

$$\text{rmse}(S) = \sqrt{|S|^{-1} \sum_{(i,u) \in S} (\hat{r}_{ui} - r_{ui})^2}$$

$S$  contains user-item  
pairs for which ratings  
are observed



# Recommender Systems

Shantanu Jain



# Content-based Filtering

Feature-based recommendation



# Item-based Features

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

# Item-based Features

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

# Item-based Features

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

# Per-user Regression

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

Learn a set of regression coefficients for each user

$$\mathbf{w}_u = \underset{\mathbf{w}}{\operatorname{argmin}} |\mathbf{r}_u - \mathbf{X}\mathbf{w}|^2 \quad \mathbf{w}_u = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{r}_u$$

Each row of  $\mathbf{X}$  contains an encoding for an item.

# Bias

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

# Bias

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9
Moonrise Kingdom	4	5	4	4	0.3	0.2

# Bias

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9
Moonrise Kingdom	4	5	4	4	0.3	0.2

*Problem:* Some movies are universally loved / hated

# Bias

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	3	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	3	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9
Moonrise Kingdom	4	3	4	4	0.3	0.2

*Problem:* Some movies are universally loved / hated  
some users are more picky than others



# Bias

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	3	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	3	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9
Moonrise Kingdom	4	3	4	4	0.3	0.2

*Problem:* Some movies are universally loved / hated  
some users are more picky than others

*Solution:* Introduce a per-movie and per-user bias

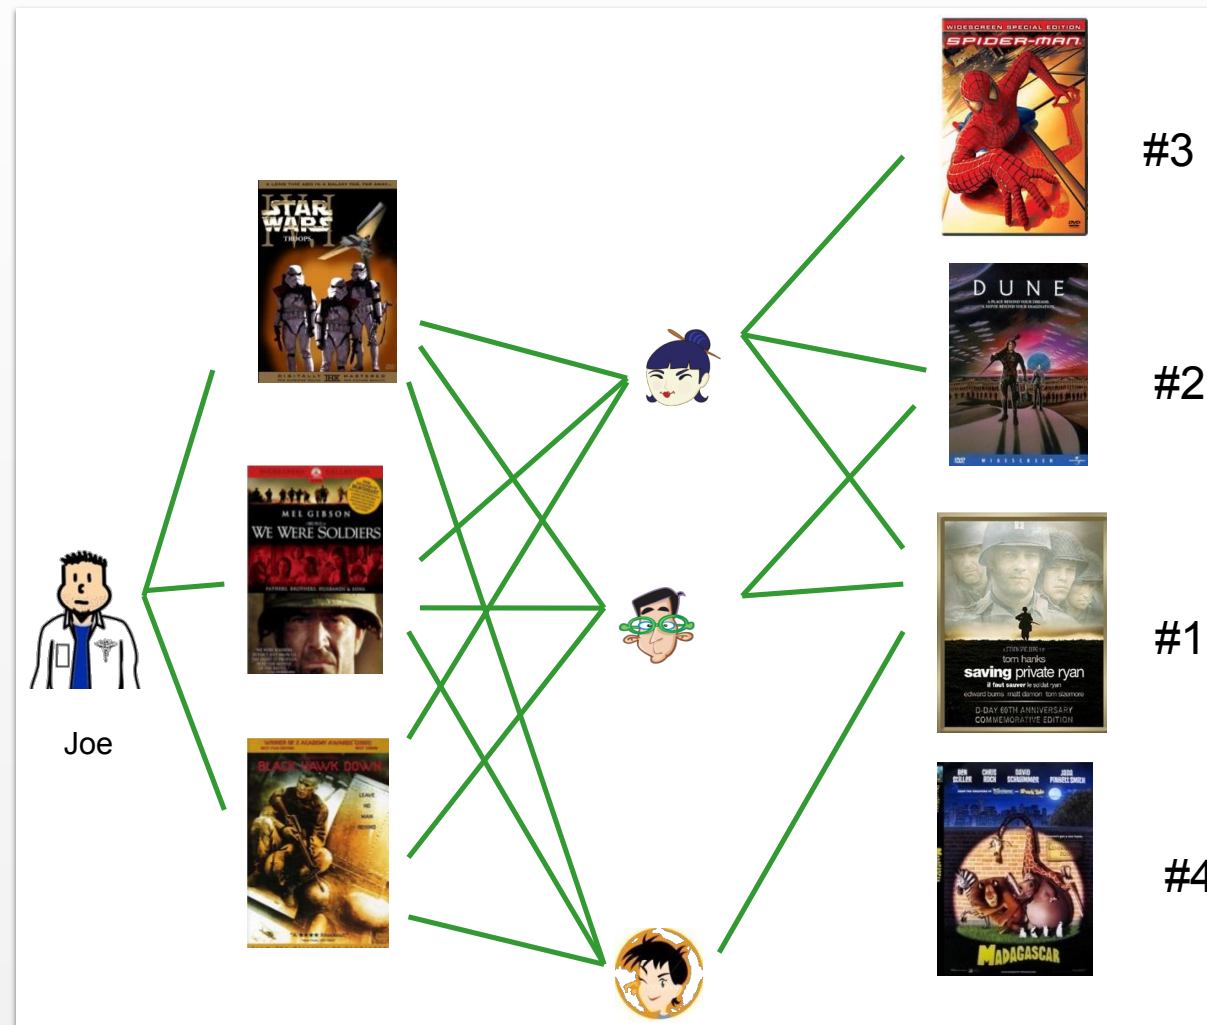
$$\hat{r}_{ui} = \mu + b_u + b_i + \mathbf{x}_i^\top \mathbf{w}_u$$



# Collaborative Filtering

Connectivity-based recommendation

# Neighborhood Based Methods



Users and items form a bipartite graph (edges are ratings)

# Neighborhood Based Methods

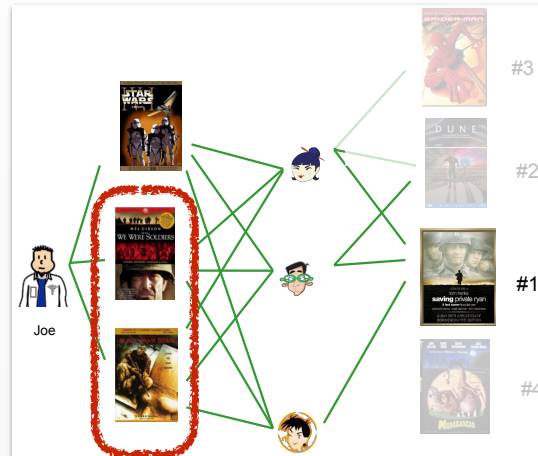
## (user, user) similarity

- predict rating based on average from k-nearest users
- good if item base changes rapidly

## (item,item) similarity

- predict rating based on average from k-nearest items
- good if user base changes rapidly

# Parzen-Window Style CF



$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in \epsilon_k(i,u)} s_{ij}(r_{uj} - b_{uj})}{\sum_{j \in \epsilon_k(i,u)} |s_{ij}|}$$

$$b_{ui} = \mu + b_u + b_i$$

- Define a similarity  $s_{ij}$  between items
- Find set  $\epsilon_k(i,u)$  of k-nearest neighbors to  $i$  that were rated by user  $u$
- Predict rating using weighted average over set
- How should we define  $s_{ij}$ ?

# Pearson Correlation Coefficient

User ratings for item  $i$ :

1	?	?	5	5	3	?	?	?	4	2	?	?	?	?	4	?	5	4	1	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

User ratings for item  $j$ :

?	?	4	2	5	?	?	1	2	5	?	?	2	?	?	3	?	?	?	5	4
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$$s_{ij} = \frac{\text{cov}(r_{.i}, r_{.j})}{\text{std}(r_{.i}) \times \text{std}(r_{.j})}$$

# (item,item) similarity

Empirical estimate of Pearson correlation coefficient

$$\hat{\rho}_{ij} = \frac{\sum_{u \in U(i,j)} (r_{ui} - b_{ui})(r_{uj} - b_{uj})}{\sqrt{\sum_{u \in U(i,j)} (r_{ui} - b_{ui})^2 \sum_{u \in U(i,j)} (r_{uj} - b_{uj})^2}}$$

$U(i, j)$ : set of users who have rated both  $i$  and  $j$

Regularize towards 0 for small support

$$s_{ij} = \frac{|U(i, j)| - 1}{|U(i, j)| - 1 + \lambda} \hat{\rho}_{ij}$$

Regularize towards baseline for small neighborhood

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in \epsilon_k(i,u)} s_{ij} (r_{uj} - b_{uj})}{\lambda + \sum_{j \in \epsilon_k(i,u)} |s_{ij}|}$$

# Similarity for binary labels

Pearson correlation not meaningful for binary labels  
(e.g. Views, Purchases, Clicks)

*Jaccard similarity*

$$s_{ij} = \frac{m_{ij}}{\alpha + m_i + m_j - m_{ij}}$$

*Observed / Expected ratio*

$$s_{ij} = \frac{\text{observed}}{\text{expected}} \approx \frac{m_{ij}}{\alpha + m_i m_j / m}$$

$m_i$  users acting on  $i$

$m_{ij}$  users acting on both  $i$  and  $j$

$m$  total number of users





# Recommender Systems

Shantanu Jain



# Matrix Factorization Methods

Learning user and item features

# Content based filtering

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	$x_1$ (romance)	$x_2$ (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

$$\hat{r}_{ui} = x_i^T w_u$$

Learn a set of regression coefficients for each user separately

$$w_u = \underset{w}{\operatorname{argmin}} |r_u - Xw|^2 \quad \text{Solution: } w_u = (X^T X)^{-1} X^T r_u$$

Each row of  $X$  contains an encoding for an item.

# Matrix Factorization (SVD)

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?
Moonrise Kingdom	4	5	4	4

This approach is referred to as SVD in the recommendation system literature, but it is not the SVD as defined in linear algebra

The predicted rating is an inner product of item and user embeddings


$$\hat{r}_{ui} = x_i^T w_u$$

In matrix form:

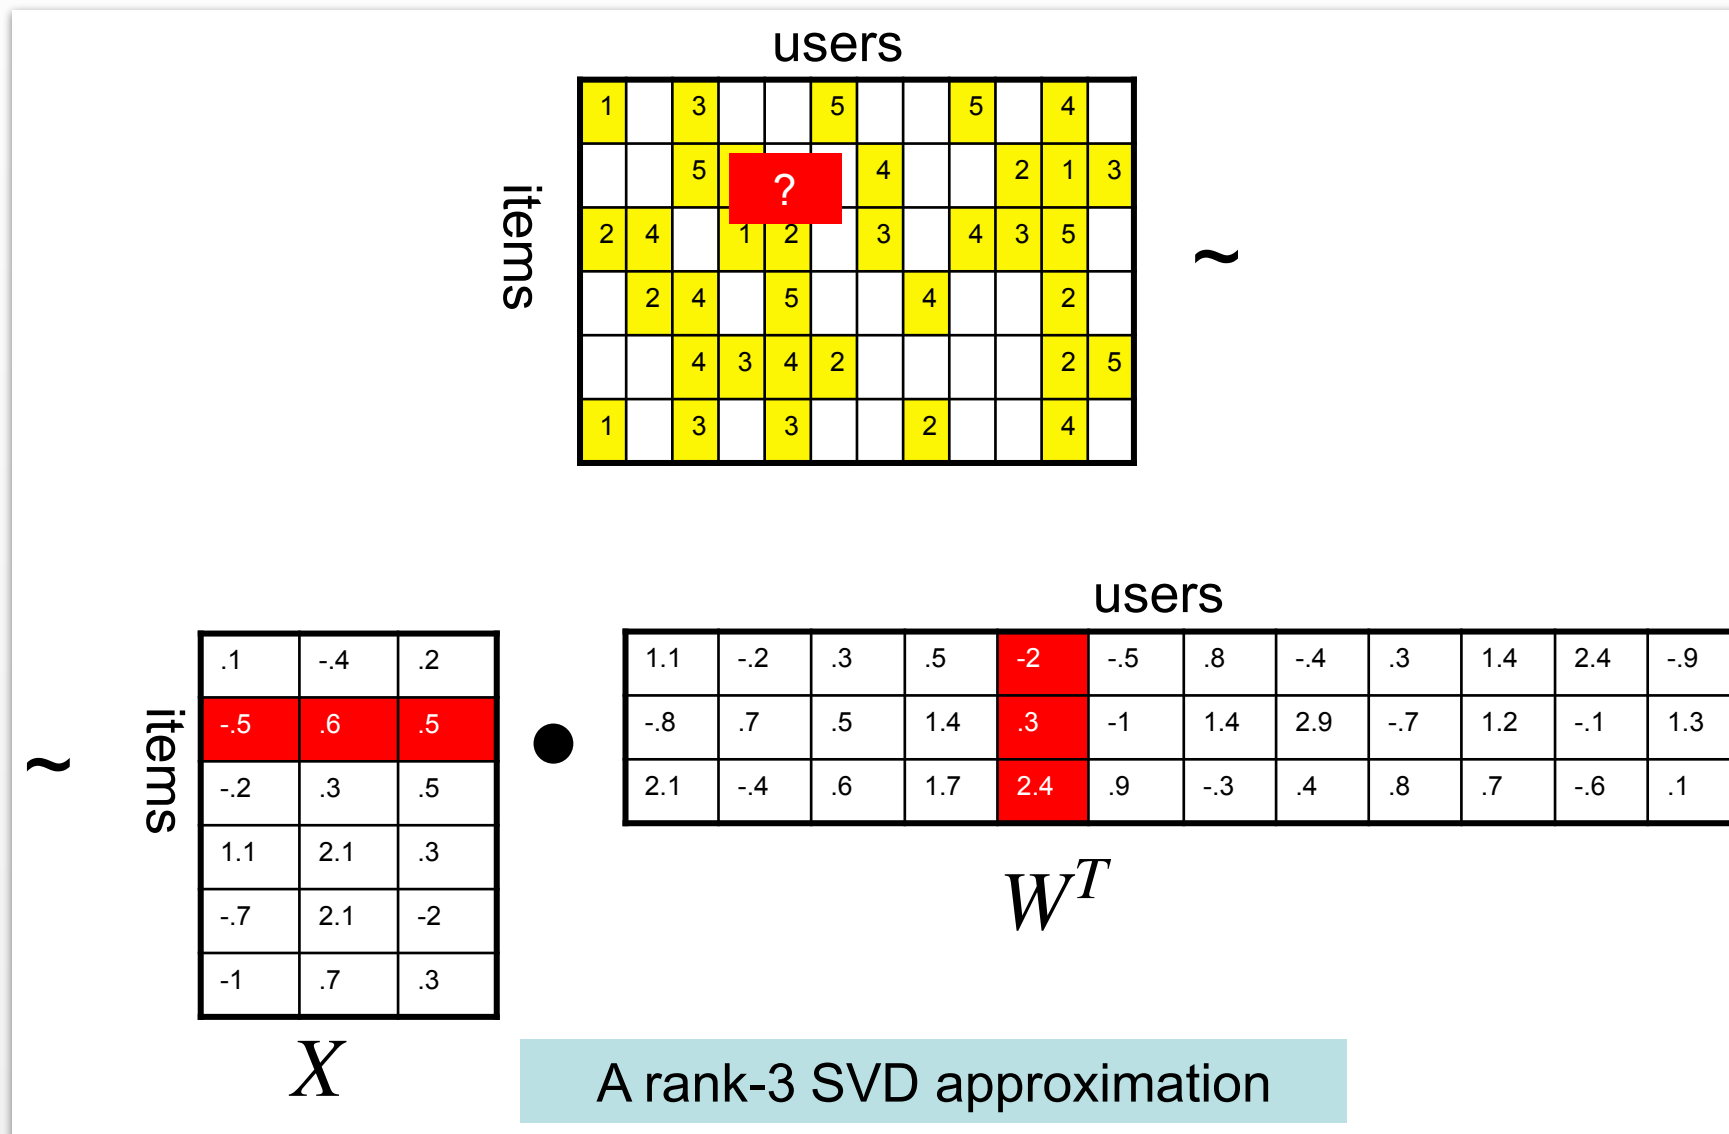
$$\hat{R} = XW^T$$

$$\hat{X}, \hat{W} = \operatorname{argmin}_{X,W} \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui})^2$$

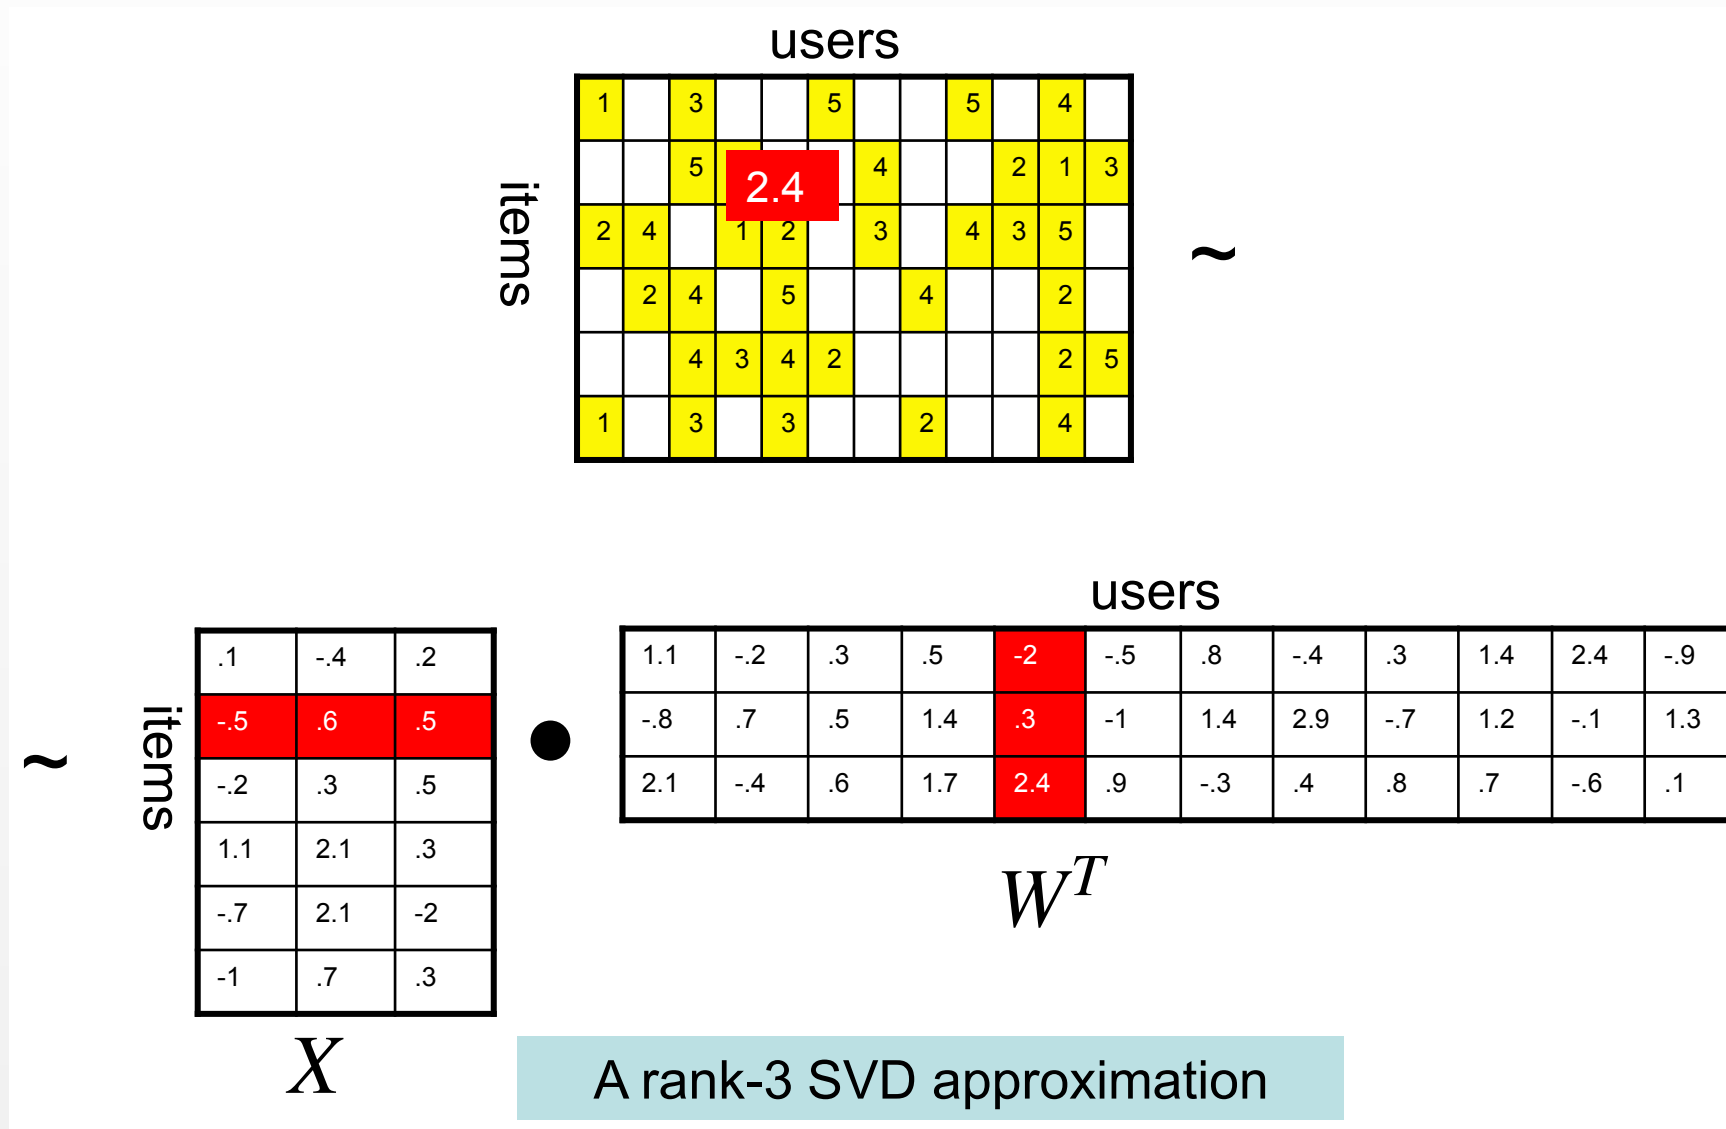
# Matrix Factorization



# Prediction



# Prediction



# Matrix Factorization

$$\hat{X}, \hat{W} = \operatorname{argmin}_{X, W} \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui})^2$$

Initialize  $X$  and  $W$  randomly

Iterate till convergence

$$1. \quad W^{(t+1)} = \operatorname{argmin}_W \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$2. \quad X^{(t+1)} = \operatorname{argmin}_X \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$



# Matrix Factorization

$$1. \quad W^{(t+1)} = \operatorname{argmin}_W \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{w}_u = (X_{S_u}^T X_{S_u})^{-1} X_{S_u}^T r_{S_u}$$

$X_{S_u}$ : contains only those items rated by  $u$ . Each row is an item encoding

$r_{S_u}$ : vector of ratings observed for user  $u$

$$2. \quad X^{(t+1)} = \operatorname{argmin}_X \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{x}_i = (W_{S_i}^T W_{S_i})^{-1} W_{S_i}^T r_{S_i}$$

$W_{S_i}$ : contains only those users that have rated item  $i$ . Each row is a user encoding

$r_{S_i}$ : vector of ratings observed for item  $i$

# SVD with missing values

1		3			5			5			4	
		5	4			4			2	1	3	
2	4		1	2		3		4	3	5		
	2	4		5			4			2		
		4	3	4	2					2	5	
1		3		3			2			4		

~

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

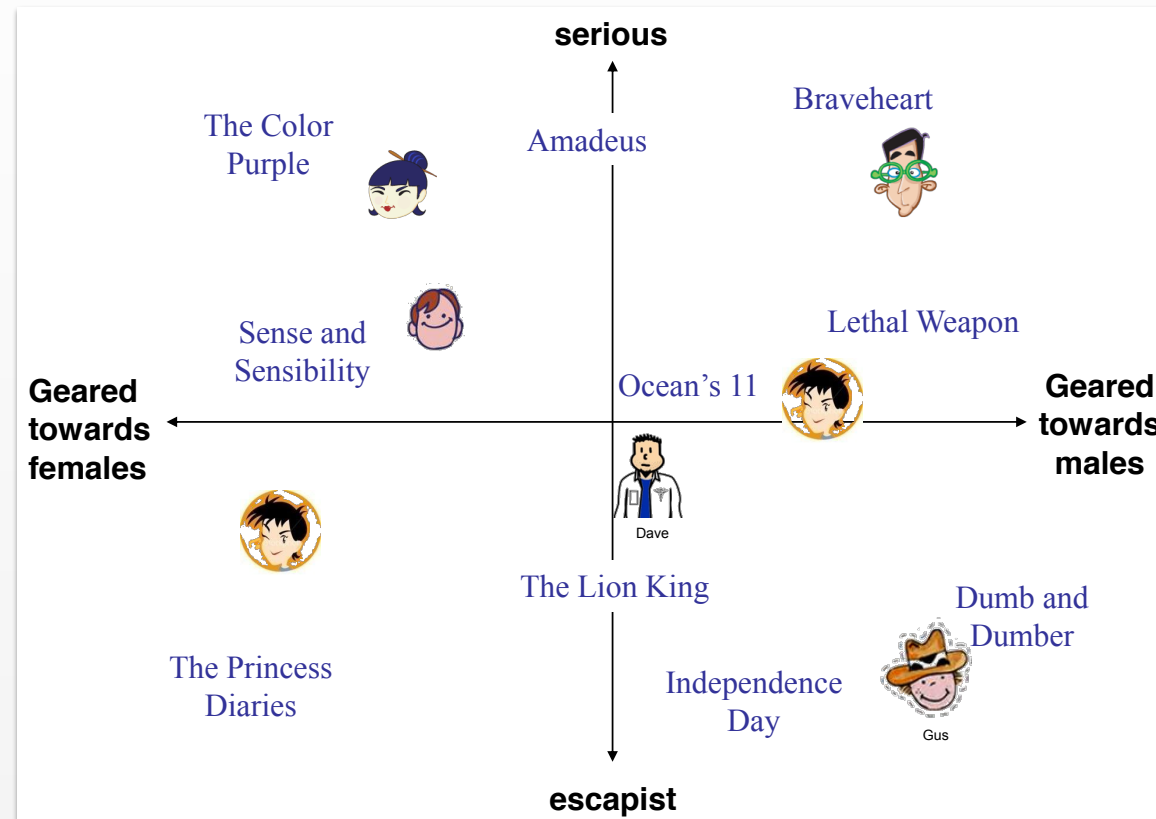
Pose as regression problem

$$\hat{X}, \hat{W} = \operatorname{argmin}_{X, W} \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui})^2 + \lambda (\|X\|_F^2 + \|W\|_F^2)$$

Regularize using Frobenius norm

$$\|A\|_F^2 = \sum_{ij} |A_{ij}|^2$$

# Need for regularization



# Matrix Factorization

$$1. \quad W^{(t+1)} = \operatorname{argmin}_W \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{w}_u = (\lambda I + X_{S_u}^T X_{S_u})^{-1} X_{S_u}^T r_u$$

$X_{S_u}$ : contains only those items rated by  $u$ . Each row is an item encoding

$r_u$ : vector of ratings observed for user  $u$

$$2. \quad X^{(t+1)} = \operatorname{argmin}_X \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{x}_i = (\lambda I + W_{S_i}^T W_{S_i})^{-1} W_{S_i}^T r_i$$

$W_{S_i}$ : contains only those users that have rated item  $i$ . Each row is a user encoding

$r_i$ : vector of ratings observed for item  $i$

# Regularized matrix Factorization

$$1. \quad W^{(t+1)} = \operatorname{argmin}_W \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{w}_u = (\lambda I + X_{S_u}^T X_{S_u})^{-1} X_{S_u}^T r_{S_u}$$

$X_{S_u}$ : contains only those items rated by  $u$ . Each row is an item encoding

$r_{S_u}$ : vector of ratings observed for user  $u$

$$2. \quad X^{(t+1)} = \operatorname{argmin}_X \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{x}_i = (\lambda I + W_{S_i}^T W_{S_i})^{-1} W_{S_i}^T r_{S_i}$$

$W_{S_i}$ : contains only those users that have rated item  $i$ . Each row is a user encoding

$r_{S_i}$ : vector of ratings observed for item  $i$

# How to pick $\lambda$ and $k$

- Split the set of observed ratings into training set and a validation set
- Train with multiple values of  $(\lambda, k)$  pairs on the training set.
- Pick the one with the best performance on the validation set.

# SVD: Adding the bias terms

$$\hat{r}_{ui} = \mu + b_u + b_i + \mathbf{x}_i^\top \mathbf{w}_u$$

$\mu$ : (global) average of all ratings across all users and movies.

$b_u$ : average difference between users  $u$ 's ratings and  $\mu$

$b_i$ : average difference between item  $i$ 's ratings and  $\mu$

$$\hat{R} = B + XW^\top$$

$$b_{ui} = \mu + b_u + b_i$$

$$B = [b_{ui}]$$

$$\hat{X}, \hat{W} = \operatorname{argmin}_{X, W} \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui})^2 + \lambda(\|X\|_F^2 + \|W\|_F^2)$$

# Regularized matrix Factorization

$$1. \quad W^{(t+1)} = \operatorname{argmin}_W \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{w}_u = (\lambda I + X_{S_u}^T X_{S_u})^{-1} X_{S_u}^T (r_{S_u} - b_{S_u})$$

$B_{S_u}$ : vector of baseline ratings  
 $\mu + b_u + b_i$  for those items  
where user  $u$  gives a rating.

$$2. \quad X^{(t+1)} = \operatorname{argmin}_X \sum_{(u,i) \in S} (r_{ui} - \hat{r}_{ui}^{(t)})^2$$

$$\hat{x}_i = (\lambda I + W_{S_i}^T W_{S_i})^{-1} W_{S_i}^T (r_{S_i} - b_{S_i})$$

$b_{S_i}$ : vector of baseline ratings  
 $\mu + b_u + b_i$  on item  $i$ , only for  
those users that indeed rate  
item  $i$ .



# Netflix Prize

## Training data

- 100 million ratings, 480,000 users, 17,770 movies
- 6 years of data: 2000-2005

## Test data


- Last few ratings of each user (2.8 million)
- Evaluation criterion: Root Mean Square Error (RMSE)

## Competition

- 2,700+ teams
- Netflix's system RMSE: 0.9514
- \$1 million prize for 10% improvement on Netflix


.

# Improvements




Do SGD, but also learn biases  $\mu$ ,  $b_u$  and  $b_i$

# Improvements




Account for fact that ratings are not missing at random.

# Improvements



$$\hat{r}_{ui} = \mu(t) + b_u(t) + b_i(t) + \mathbf{x}_i^\top \mathbf{w}_u + \mathbf{c}_u(t)^\top \mathbf{v}_i$$

# Improvements



Still pretty far from 0.8563 grand prize

# Winning Solution from BellKor

