



Mining Streams

Shantanu Jain

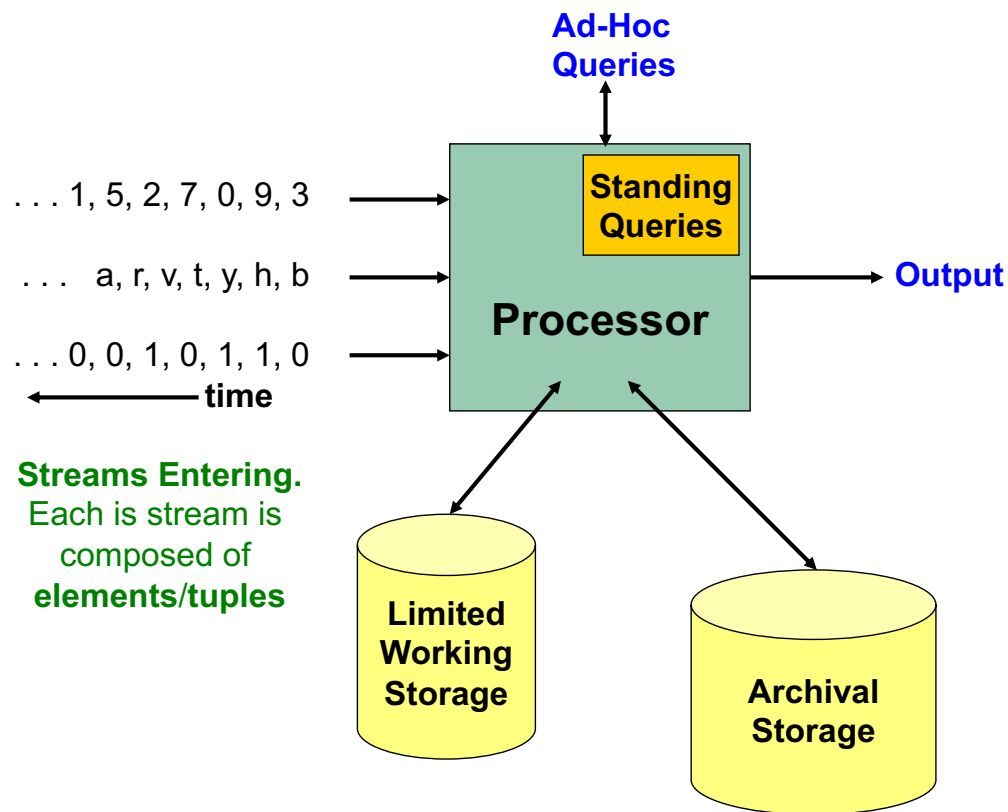
Data Streams

- In many data mining situations, we do not know what data will arrive in advance
- Stream Management is important when the input rate is controlled externally:
 - Google queries
 - Twitter or Facebook status updates
- Can think of streams as infinite and non-stationary (the distribution changes over time)

The Stream Model

- Input **elements** enter at a rapid rate, at one or more input ports (i.e., streams)
 - We often represent elements as tuples
- The system cannot store the entire stream
- Q: How do you make calculations about the stream using a limited amount of (secondary) memory?

General Stream Processing Model



- Common Types of Queries:
 - Sampling data from a stream
 - Construct a random sample
 - Queries over sliding windows
 - Number of items of type x in the last k elements of the stream
 - Filtering a data stream
 - Select elements with property x from the stream
 - Counting distinct elements
 - Number of distinct elements in last k elements of the stream
 - Estimating moments
 - Estimating frequency/surprise

Applications (1)

- Mining query streams
 - Google wants to know what queries are more frequent today than yesterday
- Mining click streams
 - Wikipedia wants to know which of its pages are getting an unusual number of hits in the past hour
- Mining social network news feeds
 - E.g., look for trending topics on Twitter, Facebook

Applications (2)

- Sensor Networks
 - Many sensors feeding into a central controller
- Telephone call records
 - Data feeds into customer bills as well as settlements between telephone companies
- IP packets monitored at a switch
 - Gather information for optimal routing
 - Detect denial-of-service attacks



Mining Streams

Shantanu Jain

Sampling from Streams

Sampling from a Data Stream

- Since we can not store the entire stream, one obvious approach is to store a **sample** (i.e. a random subset of elements)
- **Two different approaches:**
 1. Sample a **fixed fraction** of elements (say 1 in 10)
 2. Maintain a **random sample of fixed size** over a potentially infinite stream
 - At “any time” k we would like a **random sample of s elements**
 - For all time steps k , each of k elements so far should have an equal probability of being included in the s elements

Approach 1: Sampling a Fixed Proportion

- **Scenario:** Search engine query stream
 - **Stream of tuples:** (user, query, time)
 - **Answer questions such as:** How many queries from a typical user in past 30 days are repeat queries.
 - Have space to store **1/10th** of query stream
- **Naïve solution:** Random subsampling
 - Generate a random number $u \sim \text{Uniform}([0, 1))$
 - Store the query if $u < 0.1$, discard if $u \geq 0.1$

Problem with Naïve Approach

- Suppose each user issues x number of queries once and d number of queries twice (total of $x+2d$ queries)
 - True Fraction of Duplicates (unknown): $d / (x+d)$
- Naive Estimate: Keep 10% of the queries
 - Sample will contain $x / 10$ of the singleton queries and $2d / 10$ of the duplicate queries at least once
 - But only $d / 100$ pairs of duplicates
$$d/100 = 1/10 \cdot 1/10 \cdot d$$
 - Of d “duplicates” $18d / 100$ appear exactly once
$$18d / 100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$$

$$\text{Fraction in Sample} \frac{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}}{\frac{d}{100}} = \frac{d}{10x + 19d}$$

Problem with Naïve Approach

- Suppose each user issues n_x number of queries once and n_d number of queries twice (total of $n_x + 2n_d$ queries)
 - True Fraction of Duplicates (unknown): $n_d/(n_x + n_d)$
- Naive Estimate: Keep 10% of the queries
 - Sample will contain $n_x/10$ of the singleton queries and $2n_d/10$ of the duplicate queries at least once
 - But only $n_d/100$ pairs of duplicates
$$n_d/100 = 1/10 \cdot 1/10 \cdot n_d$$
 - Of d “duplicates” $18n_d / 100$ appear exactly once
$$18n_d / 100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot n_d$$

$$\text{Fraction in Sample} \frac{\frac{n_x}{10} + \frac{n_d}{100} + \frac{18n_d}{100}}{\frac{n_x}{10} + \frac{n_d}{100} + \frac{18n_d}{100}} = \frac{n_d}{10n_x + 19n_d}$$

Solution: Sample Users

Alternative Solution:

- Pick 1/10th of **users** and take all their searches in the sample
- Use a hash function to uniformly assign user names to 10 buckets

Example hash function

$$h(u) = (u \bmod 10) + 1$$

maps user id to a value in $\{1, 2, \dots, 10\}$

each user is assigned a randomly generated id

If $h(u) \leq 1$ store the tuple, else discard it.

Better Solution: Sample Keys

- Assume tuples with keys:
 - Key is some subset of each tuple's components
 - e.g., tuple is (user, search, time); key is user
 - Choice of key depends on application
- To get a sample of a / b fraction of the stream:
 - Hash each tuple's key uniformly into b buckets
 - Pick the tuple if its hash value is at most a



Hash table with b buckets, pick the tuple if its hash value is at most a .

How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first 3 buckets

Approach 2: Fixed-size Sample

- Suppose we want to maintain a random sample S of size exactly s tuples
 - E.g., main memory size constraint
 - Why? May not know length of stream in advance
- **Goal:** Ensure equal probability of inclusion
 - Suppose at time n we have seen n items
 - Each item should occur in the sample S with probability s / n

Approach 2: Fixed-size Sample

- **Algorithm (a.k.a. Reservoir Sampling)**
 - Store all the first s elements of the stream to S
 - Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s / n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random
- **Claim:** This algorithm maintains a sample S with the desired property:
 - After n elements, the sample contains each element seen so far with probability s / n

Proof: By Induction

- We prove this by induction:
 - Assume that after n elements, the sample contains each element seen so far with probability s/n
 - We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element so far with prob $s/(n+1)$
- Base case:
 - After we see $n=s$ elements the sample S has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction

- **Inductive hypothesis:** After n elements, the sample S contains each element seen so far with prob. s / n
- **Inductive step:** When element $n+1$ arrives, the probability for retention of each of the first n elements given it is already included in S is

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in } S \text{ not replaced}} = \frac{n}{n+1}$$

The unconditional probability for retention after $n+1$ steps is

$$\underbrace{\left(\frac{s}{n}\right)}_{\text{First } n \text{ Elements}} \underbrace{\left(\frac{n}{n+1}\right)}_{\substack{\text{Retained} \\ \text{after } n \text{ steps}} \quad \substack{\text{Retained in} \\ \text{step } n+1}} = \frac{s}{n+1} \quad \underbrace{\frac{s}{n+1}}_{\substack{\text{New} \\ \text{Element}} \quad \substack{\text{Element } n+1 \\ \text{not discarded}}}$$



Mining Streams

Shantanu Jain

Counting with Exponentially Decaying Windows

Sliding Windows

- One model for stream processing is to apply queries to a *window* of N most recent elements

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past Future →

Sliding Windows

- One model for stream processing is to apply queries to a *window* of N most recent elements

Stream of sales

10, 200, 100, 1000, 125, 500, 23, 72, 1250 ...

Stream of bits

1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0,

Sliding Windows

- **Difficult case:** Window size N is so large that the data cannot be stored in memory, or even on disk
 - Or, there are so many streams that windows for all cannot be stored
- **Amazon example:**
 - For every product X we keep **0/1** stream of whether that product was sold in the **n -th** transaction
 - We want answer queries, such as finding frequent items that were sold more than **s** times

Sliding Windows

- **Difficult case:** Window size N is so large that the data cannot be stored in memory, or even on disk
 - Or, there are so many streams that windows for all cannot be stored
- **Amazon example:**
 - For every product X we keep **0/1** stream of whether that product was sold in the **n -th** transaction
 - We want answer queries, such as finding frequent items that were sold more than **s** times

Exponentially Decaying Windows

- **Sliding window:** Count occurrences in last N elements

$$\sigma_t^{\text{SW}}(x) = \sum_{i=t-N}^t I[x \in A_i] \quad I[a_i \in x] = \begin{cases} 1 & x \in A_i \\ 0 & x \notin A_i \end{cases}$$

“Indicator” function:
returns 1 when query
matches, 0 when not

- **Exponentially decaying window:** Give lower “weight” to occurrences that are farther back in time

$$\sigma_t^{\text{SDW}}(x) = \sum_{i=1}^t I[x \in A_i] (1-c)^{t-i}$$

c is a small constant (e.g. 0.001) such that $(1-c)$ is close to 1, but $(1-c)^{t-i}$ decays to 0 when $t \gg i$

Exponentially Decaying Windows

- **Convenient Property:** Can compute sum at time t from sum at time $t-1$

$$\begin{aligned}\sigma_t^{\text{EDW}}(x) &= \sum_{i=1}^t I[x \in A_i] (1-c)^{t-i} \\ &= \underbrace{I[x \in A_t]}_{\text{Term for } i = t} + \underbrace{\sum_{i=1}^{t-1} I[x \in A_i] (1-c)^{t-i}}_{\text{Terms for } i < t} \\ &= I[x \in A_t] + (1-c) \sigma_{t-1}^{\text{EDW}}(x)\end{aligned}$$

Don't need to keep transactions A_1, \dots, A_t in memory,
just need to keep track of running weights $\sigma(x)$


Counting Items with Decaying Windows

$$\sigma_t^{\text{EDW}}(x) = I[x \in A_i] + (1 - c) \sigma_{t-1}^{\text{EDW}}(x)$$

- **Initialization:** Set $\sigma(x) = 0$ for all items x in some set X
- **For each time step**
 - Apply decay factor to all item weights $\sigma(x) = (1-c) \sigma(x)$
 - Increment weight for items in the current transaction A_i :
 $a \in A_i, \sigma[a] = \sigma[a]+1$

Sliding vs Exponential Windows

- Sliding and Exponential windows compute a weighted sum




- What differs is the definition of the weights

$$\text{Sliding } w_i = \begin{cases} 1 & i > t - N, \\ 0 & i \leq t - N. \end{cases} \quad \text{Exponential Decaying } w_i = (1 - c)^{t-i}$$

Sliding vs Exponential Windows

- Sliding and Exponential windows compute a weighted sum



- In a sliding window, the sum of the weights is N
- In an exponentially window, the sum is a *geometric series*

$$\lim_{t \rightarrow \infty} \sum_{i=1}^t w_i = \lim_{t \rightarrow \infty} \sum_{i=1}^t (1 - c)^{t-i} = \frac{1}{1 - (1 - c)} = \frac{1}{c}$$

We can think of $1/c$ as the “effective window size”

Extension to Itemsets

- Count (some) itemsets in an E.D.W.
 - What are currently “hot” itemsets?
 - **Problem:** Too many itemsets to keep counts of all of them in memory
- When a basket A_i comes in:
 - Multiply all counts by $(1-c)$
 - For uncounted items in A_i , create new count
 - Add 1 to count of any item in A_i and to any itemset contained in A_i that is already being counted
 - Drop counts $< \frac{1}{2}$
 - Initiate new counts (*next slide*)

Initiation of New Counts

- Start a count for an itemset $S \subseteq A_i$ if every proper subset of S had a count prior to arrival of basket B
 - **Intuitively:** If all subsets of S are being counted this means they are “frequent/hot” and thus S has a potential to be “hot”
- **Example:**
 - Start counting $S=\{i, j\}$ iff both i and j were counted prior to seeing B
 - Start counting $S=\{i, j, k\}$ iff $\{i, j\}$, $\{i, k\}$, and $\{j, k\}$ were all counted prior to seeing B



Mining Streams

Shantanu Jain

Filtering Data Streams

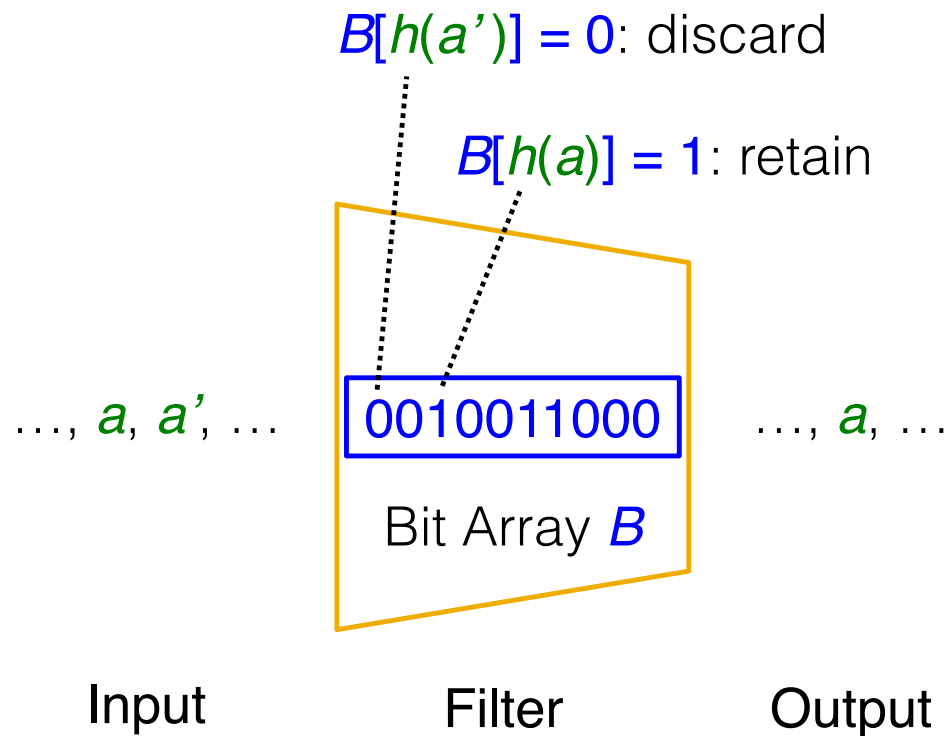
Filtering Data Streams

- Each element of data stream is a tuple
- Given a list of keys S
- Determine which tuples of stream are in S
- Obvious solution: Hash table
 - But suppose we do not have enough memory to store all of S in a hash table
 - E.g., we might be processing millions of filters on the same stream

Applications

- **Example: Email spam filtering**
 - We know 1 billion “good” email addresses
 - If an email comes from one of these, it is NOT spam
- **Publish-subscribe systems**
 - You are collecting lots of messages (news articles)
 - People express interest in certain sets of keywords
 - Determine whether each message matches user’s interest

Idea: Hash-based Filtering



- Given a set of keys S that we want to filter
 - Create a **bit array B** of n bits, initially all 0s
 - Choose a **hash function h** with range $[0, n)$
 - Hash each member of $s \in S$ to one of n buckets and set $B[h(s)] = 1$
 - For each element a , output a if $B[h(a)] == 1$
- No false negatives
- Can have false positives (hash collision)

Idea: Hash-based Filtering

Example:

S = 1 billion email addresses

B = array of 1 billion bytes (1GB)

- If the email address is in S , then it must hash to a bucket that has is set to 1, so it always gets through (*no false negatives*)
- Approximately $1/8$ of the bits are set to 1, so about $1/8^{\text{th}}$ of the addresses not in S get through (*false positives*)
 - Actually, *less* than $1/8^{\text{th}}$, because more than one address might hash to the same bit

Analysis: False Positive Rate

- More accurate analysis for the number of false positives
- **Consider:** If we throw m darts into n equally likely targets, what is the probability that a target gets at least one dart?
- **In our case:**
 - Targets = bits / buckets
 - Darts = hash values of query keys

Analysis: False Positive Rate

- We have m “darts” (hash values of items in S),
 n “targets” (bits in array)
- Assuming darts hit targets uniformly at random,
what is the probability that a target is hit?

*Probability that
all darts miss*

$$1 - \left(1 - \frac{1}{n}\right)^m = 1 - \left(\left(1 - \frac{1}{n}\right)^n\right)^{m/n} \simeq 1 - \left(\frac{1}{e}\right)^{m/n} = 1 - e^{-m/n}$$

*Probability that a
single dart misses*

*When n
is large*

Analysis: False Positive Rate

- Fraction of 1s in the array B
= probability of false positive = $1 - e^{-m/n}$
- Example: 10^9 darts, $8 \cdot 10^9$ targets
 - Fraction of 1s in $B = 1 - e^{-1/8} = 0.1175$
 - Compare with our earlier estimate: $1/8 = 0.125$

Bloom Filter

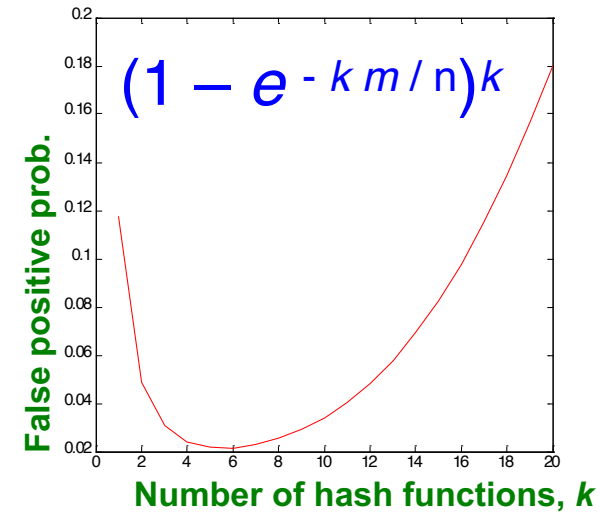
- Consider: $|S| = m$, $|B| = n$
- Use k independent hash functions h_1, \dots, h_k
- **Initialization:**
 - Set B to all 0s
 - For each $s \in S$ set $B[h_i(s)] = 1$
(for all $i = 1, \dots, k$)
- **Run-time:**
 - For each stream element with key x
 - If $B[h_i(x)] = 1$ for all $i = 1, \dots, k$
then retain x , since $x \in S$
 - Otherwise discard the element x

Bloom Filter: False Positive Rate

- What fraction of the bit vector B are 1s?
 - Throwing $k \cdot m$ “darts” at n “targets”
 - So fraction of 1s is $(1 - e^{-k m / n})$
- But we have k independent hash functions and we only let the element x through if all k hash element x to a bucket of value 1
- So, false positive probability = $(1 - e^{-k m / n})^k$

Bloom Filter: False Positive Rate

- $m = 1$ billion, $n = 8$ billion
 - $k = 1$: $(1 - e^{-1/8}) = 0.1175$
 - $k = 2$: $(1 - e^{-1/4})^2 = 0.0493$
- What happens as we keep increasing k ?
- “Optimal” value of k : $(n / m) \ln(2)$
 - In our case: Optimal $k = 8 \ln(2) = 5.54 \approx 6$
 - Error at $k = 6$: $(1 - e^{-1/6})^2 = 0.0235$



Bloom Filter: Wrap-up

- Bloom filters guarantee no false negatives, and use limited memory
 - Great for pre-processing before more expensive checks
- Suitable for hardware implementation
 - Hash function computations can be parallelized



Mining Streams

Shantanu Jain

Counting Distinct Elements

Counting Distinct Elements

- **Problem:**
 - A stream consists of a distribution over elements chosen from a set of size N
 - We would like to count the number of *distinct* elements seen so far
- **Obvious approach:**
 - Maintain set of elements seen so far
 - That is, keep a hash table of distinct elements

Applications

- How many different words are found among the Web pages being crawled at a site?
 - Unusually low or high numbers could indicate artificial pages (spam?)
- How many different Web pages does each customer request in a week?
- How many distinct products have we sold in the last week?

Using Small Storage

- Real problem: What if we do not have space to maintain the set of elements seen so far?
- Real problem: Estimate the count in an unbiased way
- Accept that the count may have a little error, but limit the probability that the error is large

Flajolet-Martin Algorithm

- Pick a hash function h that maps each of the N elements to at least $\log_2 N$ bits
- For each stream element a , let $r(a)$ be the number of trailing 0s in binary representation of $h(a)$
 - $r(a)$ = position of first **1** counting from the right
 - E.g., say $h(a) = 12$, then 12 is 1100 in binary, so $r(a) = 2$
- Record R = the maximum $r(a)$ seen
 - $R = \max_a r(a)$, over all the items a seen so far
- Estimated number of distinct elements = 2^R

Why It Works: Intuition

- Very very rough and heuristic intuition why Flajolet-Martin works:
 - $h(a)$ hashes a with equal prob. to any of N values
 - Then $h(a)$ is a sequence of $\log_2 N$ bits, where 2^{-r} fraction of all a s have a tail of r zeros
 - About 50% of a s hash to $***0$
 - About 25% of a s hash to $**00$
 - So, if we saw the longest tail of $r=2$ (i.e., item hash ending $*100$) then we have probably seen about 4 distinct items
 - So, it takes to hash about 2^r items before we see one with zero-suffix of length r

Why It Doesn't Work

- Expect value $E[2^R]$ is actually infinite
 - Probability halves when $R \rightarrow R+1$, but value doubles
- Workaround involves using many hash functions h_i and getting many samples of R_i
- How are samples R_i combined?
 - Average? What if one very large value 2^{R_i} ?
 - Median? All estimates are a power of 2
 - Pragmatic Solution:
 - Partition your samples into small groups
 - Take the median of groups
 - Then take the average of the medians



Mining Streams

Shantanu Jain

Computing Moments

Moments of a Stream

- Suppose a stream has elements chosen from a set A of with $|A| = N$ values
- Let m_i be the number of times value i occurs in the stream
- The k^{th} *moment* is $\sum_{i \in A} (m_i)^k$

Special Cases

$$\sum_{i \in A} (m_i)^k$$

- 0th moment = number of distinct elements
 - The problem just considered
- 1st moment = count of the numbers of elements (length of the stream)
 - Easy to compute
- 2nd moment = *surprise number S*
(a measure of how uneven the distribution is)

Example: Surprise Number

- Stream of length **100**
- **11** distinct values
- Item counts: 10, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9
Surprise $S = 910$
- Item counts: 90, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1
Surprise $S = 8,110$

AMS Method

- AMS method works for all moments
- Gives an unbiased estimate
- We will just concentrate on the 2nd moment S
- We pick and keep track of many variables X :
 - For each variable X , we store $X.el$ and $X.val$
 - $X.el$ corresponds to the item i
 - $X.val$ corresponds to the count of item i
 - Note this requires a count in main memory, so number of Xs is limited
- Our goal is to approximate $S = \sum_{i \in A} (m_i)^2$


One Random Variable X

- How to set $X.val$ and $X.el$?
 - Assume stream has length n (*we relax this later*)
 - Pick some random time $t < n$ to start (*any time equally likely*)
 - Set $X.el = i$, where i is the item at time t .
 - We maintain count c ($X.val = c$) of the number of i s in the stream starting from the chosen time t
- Then the AMC estimate of the 2nd moment is:

$$S = \sum_i (m_i)^2 \simeq f(X) = n(2 \cdot c - 1)$$

- Note, we can track multiple variables (X_1, X_2, \dots, X_k) to compute an average $S = \frac{1}{k} \sum_{j=1}^k f(X_j)$

AMC Estimate: Derivation



- Define c_t = number of future appearances of the item at time t
 - $c_1 = m_a, c_2 = m_a - 1, c_3 = m_b, \dots$


- Then the expected value of $f(X)=n(2c-1)$ is a sum over t

$$\mathbb{E}[f(X)] = \frac{1}{n} \sum_{t=1}^n n(2c_t - 1)$$

- We can re-arrange to obtain a sum counts c_t for each item i

$$\mathbb{E}[f(X)] = \frac{1}{n} \sum_i n \left(\underbrace{(2m_i - 1)}_{\text{First } c_t \text{ for item } i} + \underbrace{(2(m_i - 1) - 1)}_{\text{Second } c_t \text{ for item } i} + \dots + \underbrace{5 + 3 + 1}_{\text{Final } c_t \text{ for item } i} \right)$$

AMC Estimate: Derivation



- Let's rewrite the result from previous slide:

$$\begin{aligned}\mathbb{E}[f(X)] &= \frac{1}{n} \sum_i n \left((2m_i - 1) + (2(m_i - 1) - 1) + \dots + 5 + 3 + 1 \right) \\ &= \frac{1}{n} \sum_i n \sum_{j=1}^{m_i} (2j - 1) = \sum_i \sum_{j=1}^{m_i} (2j - 1)\end{aligned}$$

- Now use *triangle numbers*: $\sum_{j=1}^{m_i} j = \frac{1}{2} m_i (m_i + 1)$

$$\mathbb{E}[f(X)] = \sum_i (m_i(m_i + 1) - m_i) = \sum_i (m_i)^2 = S$$

We have now shown that $f(X)$ is an unbiased estimate of S !

Combining Samples

- **In practice:**
 - Compute $f(X) = n (2 c - 1)$ for as many variables X as you can fit in memory
 - Average them in groups
 - Take median of averages
- **Problem: Streams never end**
 - We assumed there was a number n , the number of positions in the stream
 - But real streams go on forever, so n is a variable – the number of inputs seen so far

Moments of Infinite Streams

1. The variables X have n as a factor –
keep n separately; just hold the count in X
2. Suppose we can only store k counts.
We must throw some X s out as time goes on:
 - **Objective:** Each starting time t is selected with probability k / n
 - **Solution: reservoir sampling! (*from previous video*)**
 - Choose the first k times for k variables
 - When the n^{th} element arrives ($n > k$), choose it with probability k / n
 - If you choose it, throw one of the previously stored variables X out, with equal probability