



Mining Streams

Shantanu Jain

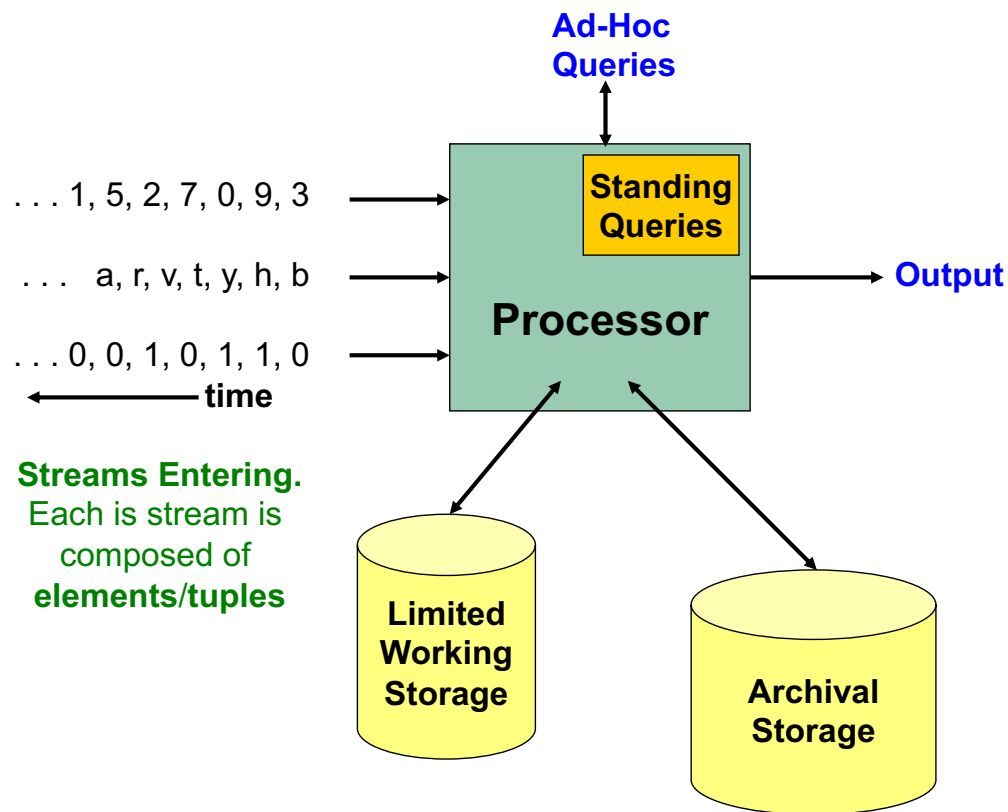
Data Streams

- In many data mining situations, we do not know what data will arrive in advance
- Stream Management is important when the input rate is controlled externally:
 - Google queries
 - Twitter or Facebook status updates
- Can think of streams as infinite and non-stationary (the distribution changes over time)

The Stream Model

- Input **elements** enter at a rapid rate, at one or more input ports (i.e., streams)
 - We often represent elements as tuples
- The system cannot store the entire stream
- **Q:** How do you make calculations about the stream using a limited amount of (secondary) memory?

General Stream Processing Model



- Common Types of Queries:
 - Sampling data from a stream
 - Construct a random sample
 - Queries over sliding windows
 - Number of items of type x in the last k elements of the stream
 - Filtering a data stream
 - Select elements with property x from the stream
 - Counting distinct elements
 - Number of distinct elements in last k elements of the stream
 - Estimating moments
 - Estimating frequency/surprise

Applications (1)

- Mining query streams
 - Google wants to know what queries are more frequent today than yesterday
- Mining click streams
 - Wikipedia wants to know which of its pages are getting an unusual number of hits in the past hour
- Mining social network news feeds
 - E.g., look for trending topics on Twitter, Facebook

Applications (2)

- Sensor Networks
 - Many sensors feeding into a central controller
- Telephone call records
 - Data feeds into customer bills as well as settlements between telephone companies
- IP packets monitored at a switch
 - Gather information for optimal routing
 - Detect denial-of-service attacks



Mining Streams

Shantanu Jain

Sampling from Streams

Sampling from a Data Stream

- Since we can not store the entire stream, one obvious approach is to store a **sample** (i.e. a random subset of elements)
- **Two different approaches:**
 1. Sample a **fixed fraction** of elements (say 1 in 10)
 2. Maintain a **random sample of fixed size** over a potentially infinite stream
 - At “any time” k we would like a **random sample of s elements**
 - For all time steps k , each of k elements so far should have an equal probability of being included in the s elements

Approach 1: Sampling a Fixed Proportion

- **Scenario:** Search engine query stream
 - **Stream of tuples:** (user, query, time)
 - **Answer questions such as:** How many queries from a typical user in past 30 days are repeat queries.
 - Have space to store **1/10th** of query stream
- **Naïve solution:** Random subsampling
 - Generate a random number $u \sim \text{Uniform}([0, 1))$
 - Store the query if $u < 0.1$, discard if $u \geq 0.1$

Problem with Naïve Approach

- Suppose each user issues x number of queries once and d number of queries twice (total of $x+2d$ queries)
 - True Fraction of Duplicates (unknown): $d / (x+d)$
- Naive Estimate: Keep 10% of the queries
 - Sample will contain $x / 10$ of the singleton queries and $2d / 10$ of the duplicate queries at least once
 - But only $d / 100$ pairs of duplicates
$$d/100 = 1/10 \cdot 1/10 \cdot d$$
 - Of d “duplicates” $18d / 100$ appear exactly once
$$18d / 100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$$

$$\text{Fraction in Sample} \frac{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}}{\frac{d}{100}} = \frac{d}{10x + 19d}$$

Problem with Naïve Approach

- Suppose each user issues n_x number of queries once and n_d number of queries twice (total of $n_x + 2n_d$ queries)
 - True Fraction of Duplicates (unknown): $n_d/(n_x + n_d)$
- Naive Estimate: Keep 10% of the queries
 - Sample will contain $n_x/10$ of the singleton queries and $2n_d/10$ of the duplicate queries at least once
 - But only $n_d/100$ pairs of duplicates
$$n_d/100 = 1/10 \cdot 1/10 \cdot n_d$$
 - Of d “duplicates” $18n_d / 100$ appear exactly once
$$18n_d / 100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot n_d$$

$$\text{Fraction in Sample} \frac{\frac{n_x}{10} + \frac{n_d}{100} + \frac{18n_d}{100}}{\frac{n_x}{10} + \frac{n_d}{100} + \frac{18n_d}{100}} = \frac{n_d}{10n_x + 19n_d}$$

Better Solution: Sample Keys

- Assume tuples with keys:
 - Key is some subset of each tuple's components
 - e.g., tuple is (user, search, time); key is user
 - Choice of key depends on application
- To get a sample of a / b fraction of the stream:
 - Hash each tuple's key uniformly into b buckets
 - Pick the tuple if its hash value is at most a



Hash table with b buckets, pick the tuple if its hash value is at most a .

How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first 3 buckets

Approach 2: Fixed-size Sample

- Suppose we want to maintain a random sample S of size exactly s tuples
 - E.g., main memory size constraint
 - Why? May not know length of stream in advance
- **Goal:** Ensure equal probability of inclusion
 - Suppose at time n we have seen n items
 - Each item should occur in the sample S with probability s / n

Approach 2: Fixed-size Sample

- **Algorithm (a.k.a. Reservoir Sampling)**
 - Store all the first s elements of the stream to S
 - Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s / n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random
- **Claim:** This algorithm maintains a sample S with the desired property:
 - After n elements, the sample contains each element seen so far with probability s / n

Proof: By Induction

- We prove this by induction:
 - Assume that after n elements, the sample contains each element seen so far with probability s/n
 - We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element so far with prob $s/(n+1)$
- Base case:
 - After we see $n=s$ elements the sample S has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction

- **Inductive hypothesis:** After n elements, the sample S contains each element seen so far with prob. s / n
- **Inductive step:** When element $n+1$ arrives, the probability for retention of each of the first n elements is

$$\left(1 - \frac{s}{n+1}\right) + \left(\frac{s}{n+1}\right) \left(\frac{s-1}{s}\right) = \frac{n}{n+1}$$

Element $n+1$
discarded

Element $n+1$
not discarded

Element in S
not replaced

The probability for retention after $n+1$ steps is

$$\begin{array}{ccc} \text{First } n & \left(\frac{s}{n}\right) \left(\frac{n}{n+1}\right) = \frac{s}{n+1} & \text{New} \\ \text{Elements} & & \text{Element} \end{array} \quad \frac{s}{n+1}$$

Retained after n steps Retained in step $n+1$ Element $n+1$ *not* discarded



Mining Streams

Shantanu Jain

Counting with Exponentially Decaying Windows

Sliding Windows

- One model for stream processing is to apply queries to a *window* of N most recent elements

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past Future →

Sliding Windows

- **Difficult case:** Window size N is so large that the data cannot be stored in memory, or even on disk
 - Or, there are so many streams that windows for all cannot be stored
- **Amazon example:**
 - For every product X we keep **0/1** stream of whether that product was sold in the **n -th** transaction
 - We want answer queries, such as finding frequent items that were sold more than **s** times

Exponentially Decaying Windows

- **Sliding window:** Count occurrences in last N elements

$$\sigma_t^{\text{SW}}(x) = \sum_{i=t-N}^t I[a_i = x] \quad I[a_i = x] = \begin{cases} 1 & a_i = x, \\ 0 & a_i \neq x. \end{cases}$$

“Indicator” function:
returns 1 when query
matches, 0 when not

- **Exponentially decaying window:** Give lower “weight” to occurrences that are farther back in time

$$\sigma_t^{\text{SDW}}(x) = \sum_{i=1}^t I[a_i = x] (1 - c)^{t-i}$$

c is a small constant (e.g. 0.001) such that $(1-c)$ is close to 1, but $(1-c)^{t-i}$ decays to 0 when $t \gg i$

Exponentially Decaying Windows

- **Convenient Property:** Can compute sum at time t from sum at time $t-1$

$$\begin{aligned}\sigma_t^{\text{EDW}}(x) &= \sum_{i=1}^t I[a_i = x] (1 - c)^{t-i} \\ &= \underbrace{I[a_t = x]}_{\text{Term for } i = t} + \underbrace{\sum_{i=1}^{t-1} I[a_i = x] (1 - c)^{t-i}}_{\text{Terms for } i < t} \\ &= I[a_t = x] + (1 - c) \sigma_{t-1}^{\text{EDW}}(x)\end{aligned}$$

Don't need to keep items $\mathbf{a}_1, \dots, \mathbf{a}_t$ in memory,
just need to keep track of running weights $\sigma(\mathbf{x})$


Counting Items with Decaying Windows

$$\sigma_t^{\text{EDW}}(x) = I[a_t = x] + (1 - c) \sigma_{t-1}^{\text{EDW}}(x)$$

- **Initialization:** Set $\sigma(x) = 0$ for all items x in some set X
- **For each new item a**
 - Apply decay factor to weights $\sigma(x) = (1-c) \sigma(x)$
 - Increment weight for current item $\sigma[a] = \sigma[a]+1$

Sliding vs Exponential Windows

- Sliding and Exponential windows compute a weighted sum




- What differs is the definition of the weights

$$\text{Sliding } w_i = \begin{cases} 1 & i > t - N, \\ 0 & i \leq t - N. \end{cases} \quad \text{Exponential Decaying } w_i = (1 - c)^{t-i}$$

Sliding vs Exponential Windows

- Sliding and Exponential windows compute a weighted sum



- In a sliding window, the sum of the weights is N
- In an exponentially window, the sum is a *geometric series*

$$\lim_{t \rightarrow \infty} \sum_{i=1}^t w_i = \lim_{t \rightarrow \infty} \sum_{i=1}^t (1 - c)^{t-i} = \frac{1}{1 - (1 - c)} = \frac{1}{c}$$

We can think of $1/c$ as the “effective window size”